

INSTITUTO TECNOLÓGICO AUTÓNOMO DE MÉXICO, ITAM
Curso: Visión por Computadora

Tarea 5

Pablo Barranco Soto 151528
Sebastián Martínez Santos 176357

Tras hacer varias búsquedas con grid search, cambiando los hiperparametros *funciones de activación, batch size, epochs, loss function, número de layers*.
Decidimos acomodar nuestros mejores modelos con respecto.

```
for a in activaciones:
    for n in number_of_layers:
        AE = Sequential()
        AE.add(Input(shape=(32, 32, 3)))
        for i in range(n):
            AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation=a))
        for i in range(n-1):
            AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation=a))
        AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))
        # Compilacion del Modelo
        for l in losses:
            for b in batch_sizes:
                for e in epochs:
                    AE.compile(optimizer='adam', loss=l, metrics = ['mse', 'acc'])
                    h = AE.fit(x=X_train, y=Y_train, batch_size=b, epochs=e, verbose=0, validation_split = 0.2)
                    print("{} {}".format( round((iteracion/324) *100,2)))
                    iteracion = iteracion + 1

                # Guardandamos los resultados
                mse_train.append(h.history['mse'][-1]) # Accuracy entrenamiento
                acc_train.append(h.history['acc'][-1]) # Accuracy validación
                loss_train.append(h.history['loss'][-1]) # Accuracy validación
```

Mejores modelos:

| Modelo | Activación | Batch Size | Epoch | Loss Function | N. Layers | Loss |
|---------------|-------------------|-------------------|--------------|--------------------------------|------------------|-------------|
| 1 | ReLu y Sigmoid | 500 | 150 | MeanSquaredError | 4 | 0.0060 |
| 2 | ReLu y Sigmoid | 500 | 50 | mean_squared_logarithmic_error | 4 | 0.0038 |
| 3 | ReLu y Sigmoid | 400 | 150 | MeanSquaredError | 4 | 0.0065 |
| 4 | ReLu y Sigmoid | 500 | 150 | MeanSquaredLogarithmicError | 4 | 0.0032 |
| 5 | Relu y Sigmoid | 800 | 200 | MeanSquaredLogarithmicError | 4 | 0.0028 |
| 6 | SeLu y Sigmoid | 500 | 50 | mean_squared_logarithmic_error | 4 | 0.0049 |
| 7 | SeLu y Sigmoid | 500 | 10 | mean_squared_logarithmic_error | 4 | 0.0077 |

Modelo 1


```
[ ] # Define model
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

AE.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-------------------|---------|
| ===== | | |
| conv2d (Conv2D) | (None, 16, 16, 8) | 224 |
| conv2d_1 (Conv2D) | (None, 8, 8, 8) | 584 |
| conv2d_transpose (Conv2DTran | (None, 16, 16, 8) | 584 |
| conv2d_transpose_1 (Conv2DTr | (None, 32, 32, 3) | 219 |
| ===== | | |
| Total params: 1,611 | | |
| Trainable params: 1,611 | | |
| Non-trainable params: 0 | | |
| ===== | | |


```
[ ] # Compile and train
AE.compile(optimizer='adam', loss=tf.keras.losses.MeanSquaredError())
h = AE.fit(x=x_train, y=y_train, batch_size=500, epochs=150, validation_split=0.2)
```



Modelo 2

```
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

#Compile
AE.compile(optimizer='Adam', loss='mean_squared_logarithmic_error')
h = AE.fit(x=x_train, y=y_train, batch_size=500, epochs=50, validation_split=0, verbose=0)
h.history['loss'][-1]
```



Modelo 3


```
[11] # Define model
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

AE.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-------------------|---------|
| ===== | | |
| conv2d (Conv2D) | (None, 16, 16, 8) | 224 |
| ----- | | |
| conv2d_1 (Conv2D) | (None, 8, 8, 8) | 584 |
| ----- | | |
| conv2d_transpose (Conv2DTran | (None, 16, 16, 8) | 584 |
| ----- | | |
| conv2d_transpose_1 (Conv2DTr | (None, 32, 32, 3) | 219 |
| ===== | | |
| Total params: 1,611 | | |
| Trainable params: 1,611 | | |
| Non-trainable params: 0 | | |
| ----- | | |

```
[22] # Compile and train
AE.compile(optimizer='adam', loss=tf.keras.losses.MeanSquaredError())
h = AE.fit(x=x_train, y=y_train, batch_size=400, epochs=150, validation_split=0.2)
```



Modelo 4


```
# Define model
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

AE.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-------------------|---------|
| conv2d (Conv2D) | (None, 16, 16, 8) | 224 |
| conv2d_1 (Conv2D) | (None, 8, 8, 8) | 584 |
| conv2d_transpose (Conv2DTran | (None, 16, 16, 8) | 584 |
| conv2d_transpose_1 (Conv2DTr | (None, 32, 32, 3) | 219 |
| Total params: 1,611 | | |
| Trainable params: 1,611 | | |
| Non-trainable params: 0 | | |

```
# Compile and train
AE.compile(optimizer='adam', loss=tf.keras.losses.MeanSquaredLogarithmicError())
h = AE.fit(x=x_train, y=y_train, batch_size=500, epochs=150, validation_split=0.2)
```



Modelo 5


```
# Define model
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='relu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

AE.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-------------------|---------|
| conv2d (Conv2D) | (None, 16, 16, 8) | 224 |
| conv2d_1 (Conv2D) | (None, 8, 8, 8) | 584 |
| conv2d_transpose (Conv2DTran | (None, 16, 16, 8) | 584 |
| conv2d_transpose_1 (Conv2DTr | (None, 32, 32, 3) | 219 |
| Total params: 1,611 | | |
| Trainable params: 1,611 | | |
| Non-trainable params: 0 | | |


```
# Compile and train
AE.compile(optimizer='adam', loss=tf.keras.losses.MeanSquaredLogarithmicError())
h = AE.fit(x=x_train, y=y_train, batch_size=800, epochs=200, validation_split=0.2)
```



Modelo 6

```
# Vemos como lucen los mejores modelos
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))


#Compile
AE.compile(optimizer='Adam', loss='mean_squared_logarithmic_error')
h = AE.fit(x=x_train, y=y_train, batch_size=500, epochs=50, validation_split=0.2, verbose = 0)
```



Modelo 7

```
AE = Sequential()
AE.add(Input(shape=(32, 32, 3)))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2D(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(8, 3, strides=(2, 2), padding='same', activation='selu'))
AE.add(Conv2DTranspose(3, 3, strides=(2, 2), padding='same', activation='sigmoid'))

#Compile
AE.compile(optimizer='Adam', loss='mean_squared_logarithmic_error')
h = AE.fit(x=x_train, y=y_train, batch_size=500, epochs=10, validation_split=0.2, verbose = 0)
```



Experiencias aprendidas:

1. Se necesita un balance para obtener buenos resultados. Es decir, más no significa mejor. Pues el aumentar el número de capas en la red neuronal o el incrementar la cantidad de “batches” y de épocas no se traduce en un mejor desempeño de la red neuronal para rellenar huecos.
2. Las funciones de activación “ReLu” y “Sigmoid” son en las que observamos mejor desempeño del algoritmo. La activación selu, difumina demasiado la imagen y a pesar de tener buenos resultados numéricos, al comparar con la imagen original, la salida de la red está muy borrosa.
3. Los resultados numéricos no siempre empatan con los mejores resultados visuales. Esto se debe a que aun cuando la pérdida *loss* fue menor para algunos experimentos no siempre se veían tan bien los resultados. Para esto lo mejor sería escoger una métrica adecuada.
4. Tuvimos dificultades para modificar el número de neuronas por capa ya que estas están sujetas a la arquitectura del autoencoder.
5. Aunque la red neuronal sea buena identificando bordes, muchas veces no tiene un buen desempeño en preservar la tonalidad de los colores de las imágenes originales.